


Clustering



Compressing Data



- Can we describe these tweets with fewer bits of information?



- Space travel tweets
- DOGE tweets
- Edgelord tweets

Compressing Data


- Can we describe these images with fewer bits of information?



- Album cover images
- Concert images
- Cat images
- Friends images


Clusters and Compression

- **Clustering is a form of data compression**
- **Many times data exists in distinct clusters**
- **If we can find these clusters, we can summarize the data in terms of the clusters**



Clusters and Generative AI


- We give the AI data for analysis
- For moderate amounts of data, we can put it all in the prompt
 - Moderate = 128,000 tokens of text
 - Moderate = 250 images
- For large amounts of data, the AI can't handle it
 - We need to compress the data for the AI
 - Clustering allows us to compress the data




Clusters and Embeddings

- Good embeddings of data map similarity into geometry
- Finding clusters then reduces to finding geometrically clustered data points

Data

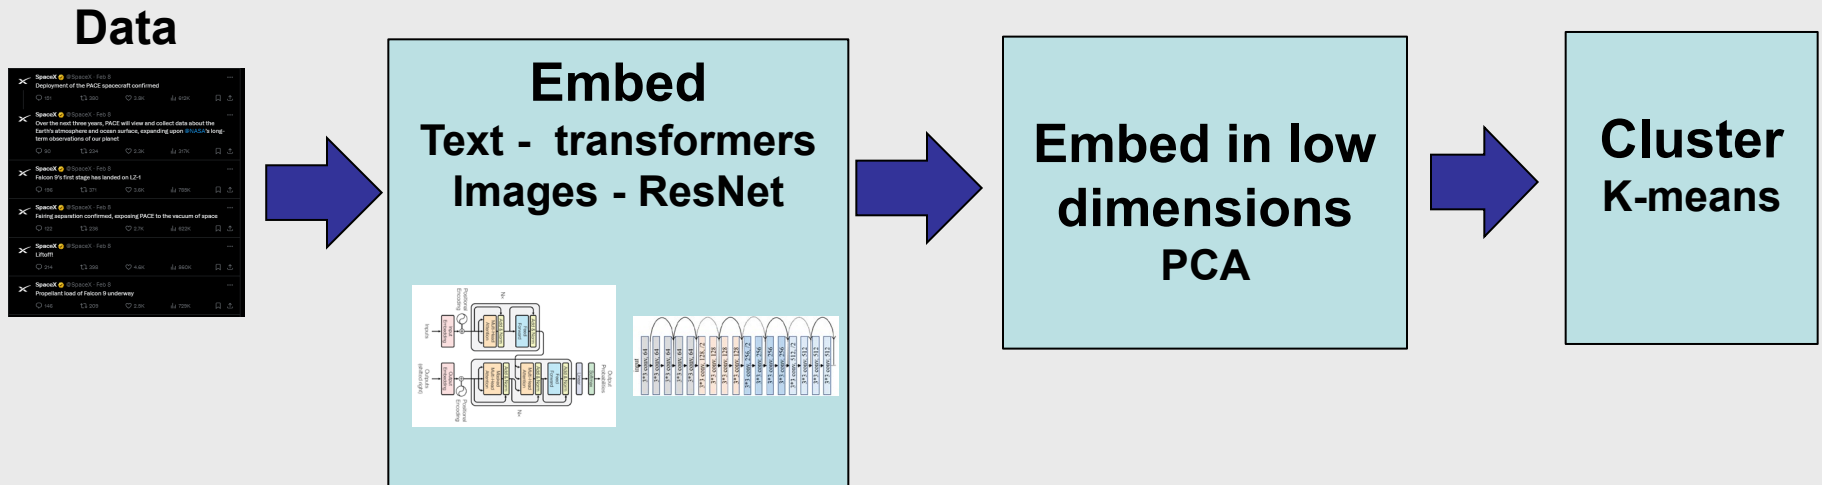



Embeddings



Clustering Pipeline


- To cluster data we will use the following steps
 - Embed with a neural network in high dimensions (transformer or ResNet)
 - Embed in low dimensions (PCA)
 - Cluster data (K-means)







PRINCIPLE COMPONENT ANALYSIS (PCA)


Embedding Data in 1 Dimension



Project
onto y





Project onto x





Embedding Data in 1 Dimension


Direction of
maximum
variance



Project
onto y




Project onto x



Direction of Maximum Variance


- If we project the data onto the direction of maximum variance, we can separate clusters
- How do we find this direction?

Direction of
maximum
variance



Covariance Matrices

- We can describe data with dimension d with a $d \times d$ covariance matrix
- This matrix encodes the direction of maximum variance




$$d = 2$$

$$\Sigma = \begin{bmatrix} E[XX] & E[XY] \\ E[YX] & E[YY] \end{bmatrix}$$

Principal Components

- The ellipse around the data is encoded in the covariance matrix
- The axes of the ellipse are the **principal components**
- The length of the axes are the standard deviations

$$\Sigma = \begin{bmatrix} 1.0 & 0.55 \\ 0.55 & 1.0 \end{bmatrix}$$



$$\Sigma = \begin{bmatrix} 1 & -1 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} 3.76 & 0 \\ 0 & 1.08 \end{bmatrix} \left(\begin{bmatrix} 1 & -1 \\ 1 & 1 \end{bmatrix} \right)^{-1}$$

Principal Component Analysis (PCA)





- **PCA lets us compute the projection of data onto the principal components (PCs) of its covariance matrix very quickly**
- **PCA was invented in 1901**
- **Has many names depending on the field you are in**
 - **Karhunen–Loève transform**
 - **Hotelling transform**
 - **eigenvalue decomposition**
 - **singular value decomposition**
 - **proper orthogonal decomposition**
 - **factor analysis**
 - **spectral decomposition**
 - **empirical orthogonal functions**

PCA Algorithm

1. Center data at zero (subtract the means)
2. Make data variance 1 in each dimension
3. Compute covariance matrix
4. Compute the principal components (PCs) with the largest variance (*fit*)
5. Project the data onto each PC to obtain the PCA embedding (*transform*)

PCA Example


- Compute covariance matrix and PC 1



- Project data to PC 1

Pros and Cons of PCA

- **Pros**
 - Extremely fast algorithm (good for big datasets)
 - Easy to interpret (each PC is a direction of high variance, probably due to some data property)
- **Cons**
 - Does not give the best embedding for some high-dimensional datasets



K-MEANS CLUSTERING


K-Means Clustering

- **K-means clustering** is a simple algorithm for clustering any kind of data (once it is embedded in a vector space)
- Invented in 1957
- Still one of the most popular ways to cluster data

K-Means Clustering Algorithm


- 1. Choose the number of clusters k**
- 2. Initialize the cluster centers randomly**
- 3. Repeat this iteration**
 - 1. Assign each data point to the cluster whose center it is closest to**
 - 2. Set the center of a cluster equal to the center of mass of the data points assigned to it**
- 4. Stop when the cluster centers stop changing**

K-Means Example



Community Detection

- An important problem in social networks is finding communities – clusters of people in a network
- Data = network structure
- Clustering algorithm = spectral clustering (K-means on spectral embeddings)




Challenges with Community Detection

- Today network data is hard to collect 😞
 - We can use Chrome plugins like TwFollow
- It's ok, we can use AI to find communities based on user (non-network) data 😊


Community Detection With Embeddings and K-Means

- User profile has a username, name, location, and bio
- Embed this data with a transformer, then use K-means to find the communities




Community Detection With AI

- User profile has a username, name, location, and bio
- We can feed all this raw data to the AI and ask it to give us the communities



AI Enhanced Clustering

- In the old days, the end goal of clustering was to assign each data point to a cluster
- Today with generative AI, we can go further and understand the clusters
 - Title
 - Description of underlying theme
 - Representative examples
- We can feed these cluster summaries to the AI for use in other tasks



Coding Session

- **Cluster tweets**
 - OpenAI transformer embeddings and K-means clustering
 - ChatGPT to describe clusters
- **Cluster images**
 - ResNet embeddings and and K-means clustering
 - ChatGPT to describe clusters
- **Find communities in a social network using no network data**
 - OpenAI transformer embeddings and K-means clustering
 - ChatGPT
 - ChatGPT to describe communities